

SparSNP — Workflow example

Gad Abraham gad.abraham@unimelb.edu.au

November 26, 2012

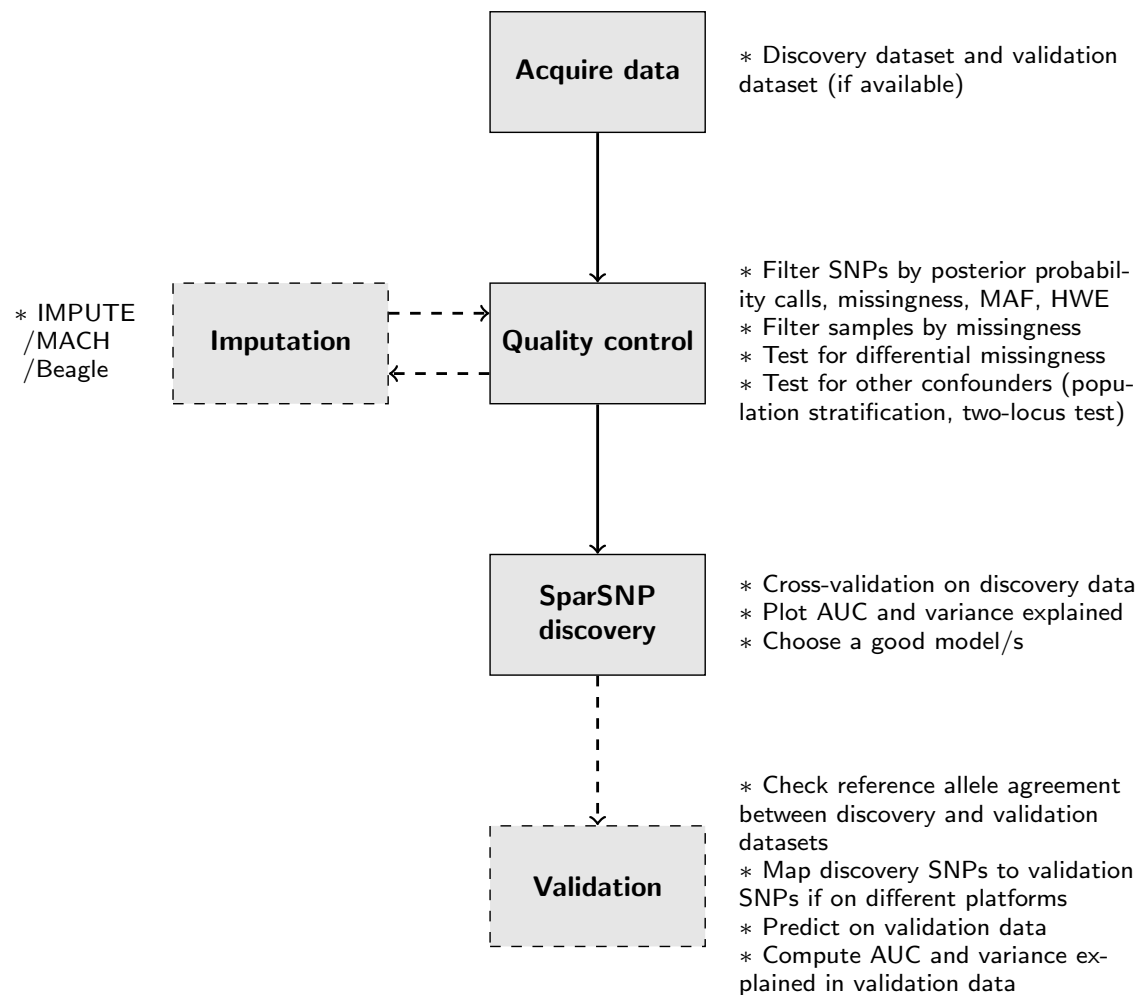


Figure 1: Analysis workflow

Here we describe a typical analysis using SparSNP. We assume that we have two datasets in PLINK BED/BIM/FAM format, named **discovery** and **validation**, on the same SNP microarray platform. The phenotype in the FAM file can be discrete (1 for controls and 2 for cases) or continuous. In the following, we assume the user is working with a Unix-like command line shell such as Bash.

1. Quality control

SparSNP implements simple imputation of genotypes such that missing genotypes are randomly assigned the value $\{0,1,2\}$ with probability of $1/3$ each, which does not introduce substantial bias to the model when the proportion of missingness is small and the genotypes are missing at random. We recommend removing samples and SNPs with high missingness, and testing for differential missingness between cases and controls.

- Remove markers with missingness ≥ 0.01 , MAF ≤ 0.01 , and Hardy-Weinberg test for equilibrium in controls of $p \leq 10^{-4}$, and samples with missingness ≥ 0.01 :

```
$ plink --bfile discovery --geno 0.01 --mind 0.01 \
--maf 0.01 --hwe 0.0001 \
--make-bed --out discovery_filtered
```
- Test for differential missingness:

```
$ plink --file discovery_filtered --missing
```


(remove SNPs as indicated by output)
- Test for sample relatedness with a threshold of $\hat{\pi} = 0.05$:

```
$ plink --file discovery_filtered --Z-genome --min 0.05
```


(remove related samples as indicated by output)
- Check for stratification using PCA, for example, using `smartpca` in Eigensoft (Price et al., 2006).
- Two locus test for detecting batch effects (Lee et al., 2010).
(remove SNPs as indicated by output)

2. SparSNP on discovery data

- Run cross-validation. By default, 10×10 -fold cross-validation will be performed, allowing up to $\min\{N, p\}$ SNPs in the model, using a case/control phenotype: for classification (case/control)

```
$ crossval.sh discovery_filtered sqrhinge
```


and for continuous outputs (linear regression)

```
$ crossval.sh discovery_filtered linear
```


The `crossval.sh` script can be modified to perform more cross-validation folds or to tune the model parameters.

You can run SparSNP on 10 cores at once using

```
$ NUMPROCS=10 crossval.sh discovery_filtered sqrhinge
```
- Plot AUC and explained phenotypic and genetic variance in the cross-validation replications, with the optional population prevalence $K = 1\%$ and heritability $h_L^2 = 50\%$:

```
$ eval.R discovery_filtered prev=0.01 h2l=0.5
```


The plots `discovery_filtered.AUC.pdf`, `discovery_filtered_VarExp.pdf`, and `discovery_filtered_GenVarExp.pdf` will be produced in the directory `discovery`. The raw AUC and explained variance data is stored in the R-data file named `discovery_filtered.RData`. The prevalence and heritability are optional. Prevalence is needed for computing explained genetic and phenotypic variance, and heritability is needed for computing explained genetic variance.

- The set of models with best predictive ability will automatically be chosen from the results, based on smoothing of the AUC or R^2 . The SNPs appearing in these models will be tabulated according to how many time they were included over all cross-validation folds. To inspect the SNPs selected by models that maximise the predictive ability:

```
head discovery/topsnps.txt
```

```
RS Counts Proportion Replications
rs2050189 60 1 60
rs2187668 60 1 60
rs9357152 60 1 60
rs7774954 60 1 60
rs3129763 58 0.966666666666667 60
...
```

The SNPs are ordered by the number of times they were included in a model with non-zero weight (Counts) out of the total number of cross-validation folds (Replications), also shown as a Proportion. SNPs at the top of the list are more stably selected by the lasso and are potentially more robust markers than SNPs at the bottom of the list.

3. Optional: Apply models to validation data

SparSNP models trained on the discovery dataset can be tested on an independent dataset, if one is available. Datasets can differ in terms of which SNPs are assayed, their ordering in the PLINK files, and their minor allele frequencies. SparSNP is completely oblivious to these differences and cannot detect them, hence we use PLINK scoring. The one caveat is that PLINK scoring does not have an intercept term, which we must add on. The process is:

- Extract a consensus model from the training data using `getmodels.R`, requesting a certain number of SNPs in the model:

```
$ getmodels.R nzreq=256
```

- Call `predict.sh` on the validation dataset which will call PLINK scoring internally:

```
$ predict.sh validation_filtered
```

Optionally, you can declare OUTDIR to specify a directory other than “predict”, and use multi-processing if you have multiple cores:

```
$ NUMPROCS=10 OUTDIR=validation_filtered_dir \
predict.sh validation_filtered
```

- Call `evalprofile.R` to compute ROC/AUC on the validation data, optionally with the prevalence to allow plotting PPV/NPV:

```
$ evalprofile.R prev=0.01
```

If you used a non-default directory for predict, then specify it:

```
$ evalprofile.R prev=0.01 outdir=validation_filtered_dir
```

The results will be saved in the directory you specified in the file `results.RData`.

4. Other post processing

- The model weights are stored in each cross-validation directory `discovery/crossvalXX/beta.csv.XX.XX` using a sparse text format `<index:weight>`, where `index` is the zero-based index of the SNP in the data (0 is the intercept), and `weight` is the model weight (a real number). The weights can be read into R (using `read.table` with `sep=":", header=FALSE`) or any other tool for visualisation or for validating the model on other datasets.

Some options that can be set by editing the file `crossval.sh` are:

- `LAMBDA2`: ℓ_2 penalty for elastic-net (default=0)
- `NFOLDS`: number of cross-validation folds (default=10)
- `NREPS`: number of cross-validation replications (default=10)
- `NZMAX`: maximum number of SNPs to allow in model (default= $\min\{N, p\}$)
- `NLAMBDA1`: number of λ penalties on the grid (default=30)
- `L1MIN`: multiplier on smallest λ used; it should be a some positive fraction such as 0.01. Setting it lower will increase the number of SNPs in the models, but will also increase computational time (default=0.001)

References

- A. L. Price, N. J. Patterson, R. M. Plenge, M. E. Weinblatt, et al. Principal components analysis corrects for stratification in genome-wide association studies. *Nat. Genet.*, 38:904–909, 2006.
- S. Lee, D. Nyholt, S. Macgregor, A. Henders, K. Zondervan, G. Montgomery, and P. Visscher. A simple and fast two-locus quality control test to detect false positives due to batch effects in genome-wide association studies. *Genet. Epidemiol.*, 34: 854–862, 2010.